

# Logging and Recovery

---

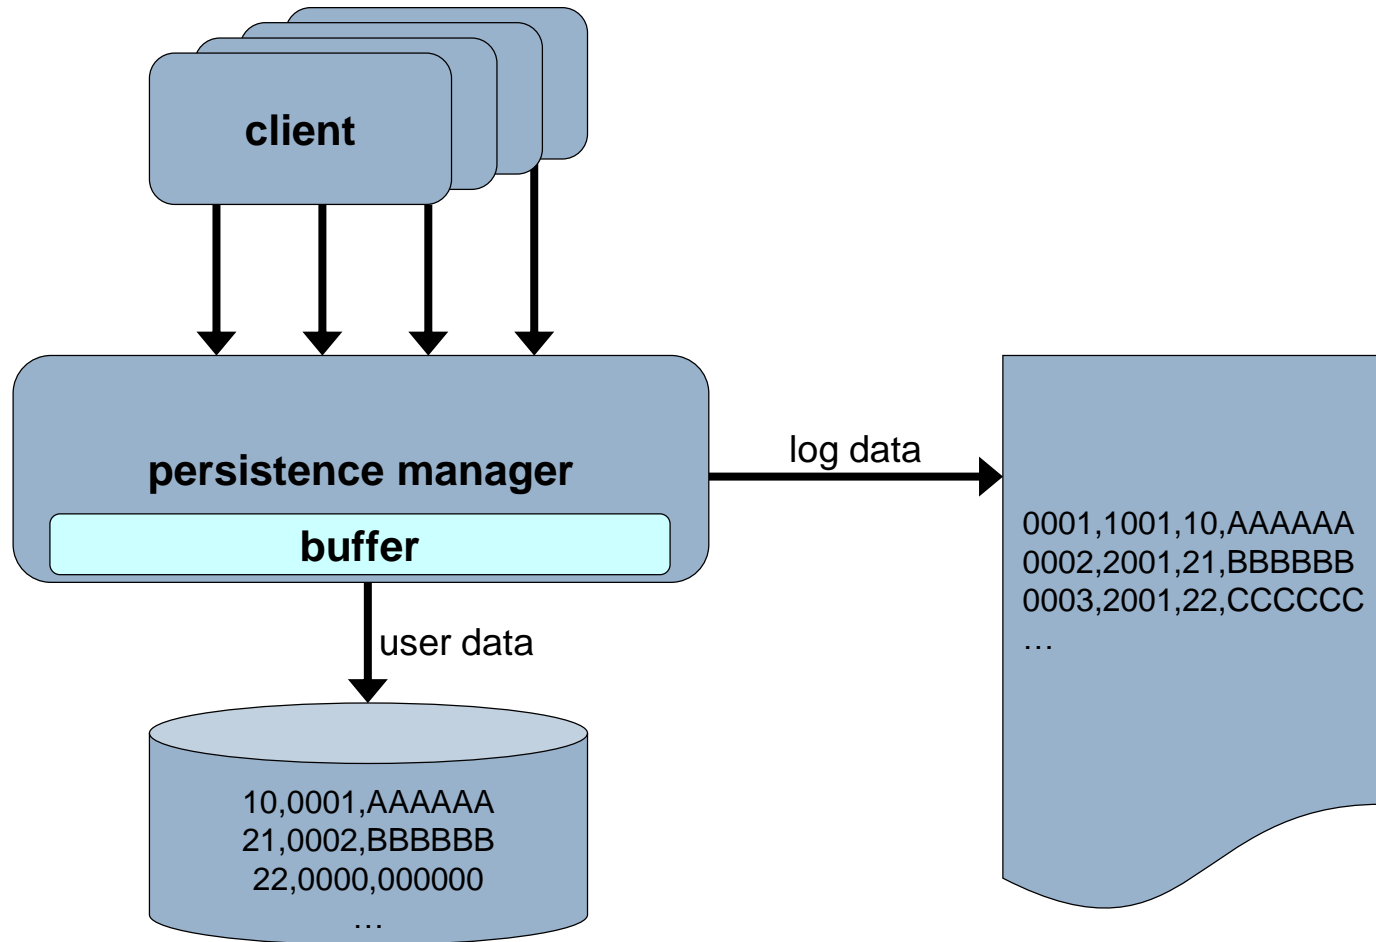
DIS Exercise Course



# Task

- implementing a simplified „database system“
- several clients access one single persistence
- buffer management
  - Non-Atomic
  - No-Steal
  - No-Force
- logging during normal operation
  - physical state logging
- crash recovery
  - redo recovery only
  - re-issuing lost modifications

# Architecture



# Persistence Manager

- persisting pages
  - in files, not in DB2
  - advice: one file per page
- page structure: [PageID, LSN, Data]
  - PageID: Page identifier
  - LSN: Log Sequence Number
  - Data: user data
- buffer: delayed writes
  - persisting data of **committed** transactions, when there are more than five pages in the buffer (No-Force)
  - check for a „full“ buffer after every write access
  - refreshing pages in the buffer is possible
  - no dirty pages are persisted (No-Steal)
  - user data are overwritten directly (Non-Atomic)
- modifying operations are logged immediately
- requirement: support of crash recovery

# Persistence Manager: Logging

- physical state logging
- granulate: Page
- new states (after-images) of modified objects are written to the log file
- record types
  - BOT and commit record
  - modification records (after-images)
- structure of log entries: [LSN, TAID, PageID, Redo]
  - LSN: Log Sequence Number (monotonically ascending)
  - TAID: transaction ID
  - PageID: Page ID
  - Redo: information required for redo

# Clients

- 5 clients of the same kind
  - concurrent threads with ClientIDs 1..5
- access to persistence manager
  - persistence manager as Singleton: a single instance
  - clients write pages like this:  
`beginTransaction() write() write()... commit()`
  - no concurrent access on the same object → no locks required  
(Client 1: pages 10..19, Client 2: pages 20..29 etc.)
  - TAID is assigned by the persistence manager

# Crash Recovery

- a record in the DB is either
  - up-to-date
  - or stale (noforce) → redo
- analysis phase
  - read the entire log sequentially
  - determine winner TAs
- redo phase
  - read the entire log sequentially
  - selective redo (redo winners)
- recovery progress is logged by updating the LSNs of the completed redo steps in the corresponding pages

# Crash Recovery

- redo is only required, if  
page LSN < LSN of redo log record
- page LSN is updated on redo (grows monotonically)

```
if LSN(LS) > LSN(Page) then  
    redo (modification from LS);  
    LSN(Page) := LSN(LS);  
end;
```



# Thread-Safe Singleton

```
class Singleton {  
  
    static final private Singleton singleton;  
    static {  
        try {  
            singleton = new Singleton();  
        }  
        catch (Throwable e) {  
            throw new RuntimeException(e.getMessage());  
        }  
    }  
  
    private Singleton() {}  
  
    static public Singleton getInstance() {  
        return singleton;  
    }  
}
```

# Creating Threads (1/2)

```
public class HelloThread extends Thread {  
    String name;  
  
    public HelloThread(String name){  
        this.name = name;  
    }  
  
    public void run() {  
        while(true){  
            System.out.println("Hello from " + name);  
            try{  
                Thread.sleep(2000);  
            }catch (InterruptedException e){  
                return;  
            }  
        }  
    }  
}
```

# Creating Threads (2/2)

```
public class ThreadCreator {  
  
    public static void main(String[] args) {  
        Thread t1 = new HelloThread("Thread1");  
        Thread t2 = new HelloThread("Thread2");  
  
        t1.start();  
        t2.start();  
  
        try {  
            Thread.sleep(10000);  
        } catch (InterruptedException e) {  
        }  
  
        t1.interrupt();  
        t2.interrupt();  
    }  
}
```

# Mutual Exclusion

```
public class SynchronizedCounter {  
    private int c = 0;  
  
    public synchronized void increment() {  
        c++;  
    }  
  
    public synchronized void decrement() {  
        c--;  
    }  
  
    public synchronized int value() {  
        return c;  
    }  
}
```